

Data Preprocessing & Model Tuning

Machine Learning Workshop – Day 2

Erdener Emin Eker

May 11, 2025

TEDU AI Data Science & AFIT Workshop

What We'll Cover Today

Main Topics:

- Cleaning data: missing values and outliers
- Feature engineering: encoding and transforming features
- Scaling: MinMax, Standard, and Robust scalers
- Model evaluation setup: train/test split and cross-validation
- Hyperparameter tuning: Grid Search and Random Search
- Building robust ML Pipelines

Goal: Prepare high-quality inputs for machine learning models and tune them effectively.

Why Preprocessing Matters

- Real-world datasets are noisy, incomplete, and inconsistent.
- ML models are sensitive to the structure and scale of input data.
- Poor preprocessing can lead to:
 - Biased or misleading predictions
 - Poor generalization to unseen data
 - Model failures in production
- Preprocessing is not a side task — it's core to ML success.

Dimensions of Data Quality

Validity: Ensures data conforms to defined rules and formats (e.g. correct type, range limits).

Accuracy: Measures how well data reflects true values (e.g. valid street address may still not exist).

Completeness: Are any required fields missing? Are we covering all necessary variables?

Consistency: No contradictions within or across datasets (e.g. a child marked as "married").

Uniformity: Standardization of formats and units (e.g. kg vs lbs, date formats).

Practical Steps in Data Cleaning

Step 1: Remove Duplicate or Irrelevant Observations Eliminate repeated rows and irrelevant entries that don't serve the analysis goal.

Step 2: Fix Structural Errors Correct typos, inconsistent capitalizations, and malformed categories (e.g. "N/A" vs "Not Applicable").

Step 3: Filter Unwanted Outliers Assess validity of extreme values — either treat, flag, or exclude them if inappropriate.

Step 4: Handle Missing Data Choose from dropping, imputing, or redesigning how the missing data is handled. Consider the reason for missingness.

Step 5: Validate and QA Ask: Does the data make sense? Does it support or contradict your assumptions?

A Real-World Dataset is Not a Clean CSV

Common Issues:

- Missing values and duplicates
- Mixed data types
- Outliers and rare categories
- Inconsistent encoding

| color | director_name | duration | gross | movie_title | language | country | budget | title_year | imdb_score |
|-------|---------------------|----------|-----------|---|----------|-------------|-----------|------------|------------|
| Color | Martin Scorsese | 240 | 116866727 | The Wolf of Wall Street | English | USA | 100000000 | 2013 | 8.2 |
| Color | Shane Black | 195 | 408992272 | Iron Man 3 | English | USA | 200000000 | 2013 | 7.2 |
| color | Quentin Tarantino | 187 | 54116191 | The Hateful Eight | English | USA | 440000000 | 2015 | 7.9 |
| Color | Kenneth Lonergan | 186 | 46495 | Margaret | English | usa | 14000000 | 2011 | 6.5 |
| Color | Peter Jackson | 186 | 258355354 | The Hobbit: The Desolation of Smaug | English | USA | 225000000 | 2013 | 7.9 |
| | N/A | 183 | 330249062 | Batman v Superman: Dawn of Justice | English | USA | 250000000 | 2017 | 6.9 |
| Color | Peter Jackson | -50 | 303001229 | The Hobbit: An Unexpected Journey | English | USA | 180000000 | 2012 | 7.9 |
| Color | Edward Hall | 180 | | Restless | English | UK | | 2012 | 7.2 |
| Color | Joss Whedon | 173 | 623279547 | The Avengers | English | USA | 220000000 | 2012 | 8.1 |
| Color | Joss Whedon | 173 | 623279547 | The Avengers | English | USA | 220000000 | 2012 | 8.1 |
| | Tom Tykwer | 172 | 27058580 | Cloud Atlas | English | Germany | 102000000 | 2012 | -7.5 |
| Color | Null | 158 | 102515793 | The Girl with the Dragon Tattoo | English | USA | 90000000 | 2011 | 7.8 |
| Color | Christopher Spencer | 170 | 59696176 | Son of God | English | USA | 22000000 | 2014 | 5.6 |
| Color | Peter Jackson | 164 | 255108370 | The Hobbit: The Battle of the Five Armies | English | New Zealand | 250000000 | 2014 | 7.5 |
| Color | Tom Hooper | 158 | 148775460 | Les Misérables | English | USA | 61000000 | 2012 | 7.6 |
| Color | Tom Hooper | 158 | 148775460 | Les Misérables | English | USA | 61000000 | 2012 | 7.6 |

(Example: A sample messy dataset)

Types of Missing Data

Not all missing data is the same:

- **MCAR** – Missing Completely at Random → e.g., survey skipped by mistake
- **MAR** – Missing at Random → e.g., income missing, but depends on education
- **MNAR** – Missing Not at Random → e.g., people with high income skip income question

Why it matters: The type of missingness affects how we handle it and the bias introduced.

Strategies for Handling Missing Data

- **Remove:**

- Drop rows/columns with missing values (if few)

- **Simple Imputation:**

- Mean, median, or mode substitution


- **Advanced Imputation:**

- KNN imputation
- Regression imputation
- Multivariate Imputation by Chained Equations (MICE)

Caveat: Imputation adds assumptions. Choose based on context and model sensitivity.

Visualizing Missing Values

- Detecting missing data is the first step in cleaning
- Useful tools in Python:
 - `'df.isnull().sum()'`
 - `'missingno.matrix(df)'`
 - `'sns.heatmap(df.isnull())'`



(Visualizing missing patterns helps identify data issues)

Comparing Imputation Techniques

Each imputation method has trade-offs:

- **Mean Imputation:**

- Fills missing values with the variable's mean.
- Simple and fast, but sensitive to outliers.
- Shrinks variance and can distort correlations.

- **Median Imputation:**

- Uses the median value — more robust to outliers.
- Preserves central tendency but still reduces variability.
- Ideal when data is skewed.


- **K-Nearest Neighbors (KNN) Imputation:**

- Estimates missing values based on similar observations.
- Preserves multivariate relationships better.
- Slower and more complex; sensitive to scaling.

Insight: Choose the method based on the missingness mechanism, variable distribution, and your model's sensitivity to distortions.

Impact on Distributions


- Imputation isn't just about filling blanks
- It can significantly affect downstream model behavior
- Visualize before/after to ensure you're not distorting key signals



(Compare variable histogram before vs. after filling missing values)

What is an Outlier?

- An outlier is an observation that lies an abnormal distance from other values.
- Causes can include:
 - Measurement error
 - Data entry mistake
 - Rare, but valid event
- Outliers can distort:
 - Mean and standard deviation
 - Model coefficients (especially in linear models)
- Always inspect context before removal.




(Example: boxplot highlighting an outlier)

Outlier Detection Methods

Univariate methods:

- **Z-score:** Identifies values far from the mean. Rule of thumb: beyond 3 standard deviations.
- **Interquartile Range (IQR):** Flags values outside $1.5 \times \text{IQR}$ from Q1 or Q3. Robust to non-normal data.



Tip: Use multiple methods to cross-check results.

Visual Outlier Detection

- Visualization helps identify patterns that statistics may miss.
- Common tools:
 - Boxplots – highlight extreme values in distributions
 - Scatterplots – detect bivariate anomalies
 - Histograms – reveal skewness or long tails
- Combine visual inspection with formal methods for a comprehensive approach.

Outlier handling should be tailored to your problem, not done automatically.

What is Feature Engineering?

- The process of creating, transforming, or selecting variables to improve model performance.
- Involves both domain knowledge and data understanding.
- Better features often outperform more complex models.
- Goals:
 - Improve predictive power
 - Make data more interpretable or structured
 - Handle raw or unstructured data effectively

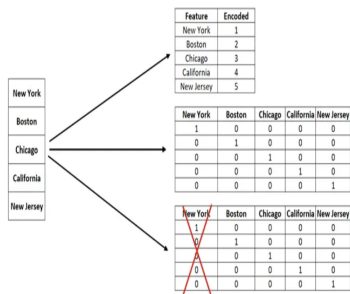
Encoding Categorical Variables

Why encode?

- Categorical features must be converted into numerical format.

Common Techniques:

- **Label Encoding:** Assigns a unique integer to each category. Implies ordinal relationship.
- **One-Hot Encoding:** Creates a new binary column for each category. No order is assumed.
- **Dummy Encoding:** Creates binary columns like One-Hot Encoding but drops one category to avoid multicollinearity. Typically used in linear models.



(Example: different encoding strategies for "Cities")

When to Use Which Encoding?

- **Tree-based models (e.g., Random Forest, XGBoost):**
 - Can handle label-encoded variables fairly well.
 - Prefer label or frequency encoding to reduce dimensionality.
- **Linear models (e.g., Logistic Regression, SVM):**
 - Require one-hot encoding for categorical variables.
 - Otherwise, model assumes incorrect linear relationship between categories.
- **High-cardinality categorical variables:**
 - One-hot encoding may lead to sparse, high-dimensional data.
 - Consider grouping infrequent categories or using frequency/target encoding.

Feature Transformation

- Transformations can help linearize relationships and stabilize variance.
- Common transformations:
 - **Log Transform:** Compresses long right tails.
 - **Square Root:** Handles moderate skewness.
 - **Box-Cox or Yeo-Johnson:** Data-driven, preserves normality.
 - **Polynomial Features:** Add non-linearity to linear models.
- Apply transformations with care to maintain interpretability.

Date and Time Feature Engineering

- Raw date variables are rarely useful in original form.
- Extract useful components:
 - Day, month, year, hour
 - Day of week, weekend flag
 - Time since event or rolling time windows
- Important for seasonality, trend detection, and time-aware models.
- Avoid leakage: do not extract future-based features in predictive models.

Why Scaling is Necessary

- Many machine learning models are sensitive to the scale of features.
- Unscaled features can dominate the learning process, especially in:
 - Distance-based models: k-NN, SVM
 - Gradient-based models: Logistic Regression, Neural Networks
- Scaling helps ensure that:
 - All features contribute equally
 - Optimization converges more smoothly
- Tree-based models (e.g., Random Forest, XGBoost) are generally scale-invariant.

Types of Feature Scaling Methods

Standard techniques:

- **StandardScaler:**

- Centers features around zero with unit variance.
- Assumes approximately normal distribution.

- **MinMaxScaler:**

- Scales features to a fixed range, usually $[0, 1]$.
- Sensitive to outliers.

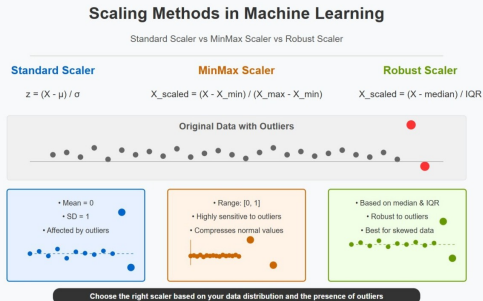
- **RobustScaler:**

- Uses median and IQR, making it robust to outliers.
- Good for skewed distributions.

Choose the scaler based on your data distribution and model type.

Visualizing the Impact of Scaling

- Scaling changes the shape and range of data.
- Important to visualize before and after to detect unintended distortions.
- Helps debug downstream model performance issues.

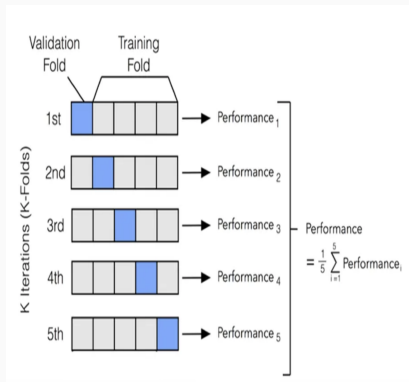


Train/Test Split

- Splitting data is essential to measure out-of-sample performance.
- Typical split: 70–80
- Important considerations:
 - **Stratification:** Ensures class balance in classification problems.
 - **Random seed:** Enables reproducibility.
 - **No leakage:** Do not scale or impute using test data.
- `train_test_split()` from `sklearn.model_selection` handles this efficiently.

Cross-Validation

- To evaluate model performance on multiple subsets.
- Common types:
 - **k-Fold CV:** Split into k equal folds; each used once as validation.
 - **Stratified k-Fold:** Maintains class balance in each fold.
 - **TimeSeriesSplit:** Maintains temporal order for forecasting.
- Reduces the risk of overfitting
- Essential for reliable model validation.



(Illustration: how data is split in k-Fold CV)

Common Pitfalls Without Cross-Validation

- Evaluating performance on a single test set can be misleading.
- Models might perform well on the split by chance (data leakage or overfitting).
- Hyperparameters tuned on a single test set may not generalize.
- Cross-validation provides a more stable estimate of model quality.
- Essential for small datasets where a single test split isn't representative.

Good models generalize — cross-validation helps prove it.

What is Hyperparameter Tuning?

- Hyperparameters are configuration settings not learned during model training.
- Examples:
 - Number of neighbors in k-NN
 - Maximum tree depth in decision trees
 - Regularization strength in linear models
- Tuning these values can significantly affect model performance.
- The goal is to find the combination that yields the best validation performance.

Grid Search vs Random Search


Grid Search:

- Exhaustive search over a defined parameter grid.
- Guarantees evaluation of all combinations.
- Can be computationally expensive.

Random Search:

- Samples parameter combinations randomly.
- More efficient with large hyperparameter spaces.
- Often finds good solutions faster than grid search.

Both require cross-validation to avoid overfitting to the validation set.



(Comparison of search coverage in parameter space)

Choosing the Right Evaluation Metric

- **Classification:**

- Accuracy — simple but misleading on imbalanced data
- Precision, Recall, F1 — for skewed classes
- ROC-AUC — performance across thresholds

- **Regression:**

- Mean Squared Error (MSE)
- Mean Absolute Error (MAE)
- R^2 Score

Tuning without the right metric can optimize the wrong outcome.

Always align your metric with the business or research goal.

| | | Predicted Class | | |
|--------------|----------|-------------------------------------|---|---|
| | | Positive | Negative | |
| Actual Class | Positive | True Positive (TP) | False Negative (FN) Type II Error | Sensitivity $\frac{TP}{(TP + FN)}$ |
| | Negative | False Positive (FP) Type I Error | True Negative (TN) | Specificity $\frac{TN}{(TN + FP)}$ |
| | | Precision $\frac{TP}{(TP + FP)}$ | Negative Predictive Value $\frac{TN}{(TN + FN)}$ | Accuracy $\frac{TP + TN}{(TP + TN + FP + FN)}$ |

(Confusion Matrix: TP, FP, FN, TN)

Why Use Pipelines?

- Pipelines combine preprocessing and modeling steps into a single object.
- Advantages:
 - Reduces code repetition and clutter
 - Prevents data leakage during preprocessing
 - Ensures consistent transformation across training and test sets
 - Enables easy deployment and cross-validation
- Provided by `sklearn.pipeline.Pipeline`

Think of a pipeline as an ML assembly line.

Typical Pipeline Structure

A simple example:

- `SimpleImputer()` → `StandardScaler()` → `LogisticRegression()`
- Encapsulated as:

```
Pipeline(steps=[  
    ("impute", SimpleImputer()),  
    ("scale", StandardScaler()),  
    ("model", LogisticRegression())  
])
```

- Pipelines can include encoding, scaling, feature selection, modeling, etc.
- For heterogeneous columns, use `ColumnTransformer`.

Cross-Validation and Tuning with Pipelines

- Pipelines integrate seamlessly with `GridSearchCV` and `RandomizedSearchCV`.
- Example:

```
param_grid = {  
    "model__C": [0.01, 0.1, 1, 10]  
}
```

```
grid = GridSearchCV(pipeline, param_grid, cv=5)
```

- Hyperparameters can be tuned across all steps using double underscores.
- Allows reproducible and clean ML workflows from start to finish.

Summary of Best Practices

- Always inspect and handle missing values carefully.
- Detect and justify treatment of outliers — don't remove blindly.
- Choose encoding and scaling based on model type and data characteristics.
- Use train/test split and cross-validation to avoid overfitting.
- Tune models using appropriate metrics and CV strategies.
- Wrap preprocessing and modeling in pipelines to ensure reproducibility.

Good preprocessing and validation are the foundation of any reliable ML workflow.

→ Let's switch to Colab!